

the belief — structure should help more as tasks get longer

SDD invokes its specification chain **2.2× more often** on long-horizon tasks

2.30 vs 1.06 Skill calls per task ($p = 0.093$) · $\rho = +0.379$ vs log expert time ($p = 0.056$)



what the chain then does — budget and context go into authoring markdown specifications;
zero recognised verification commands across 20 ceiling trials



what it costs — the largest partial-credit deficit anywhere in the data:
-0.182 against baseline on long-horizon tasks, versus **-0.080** on short ones



the paradox — the agent acts on the length hypothesis exactly where the hypothesis predicts a payoff, and that is where it does worst