



Both toolkits engage more where the hypothesis says they should pay off



Mann-Whitney $p = 0.469$ · rho vs log expert time $+0.197$ ($p = 0.335$)



Mann-Whitney $p = 0.093$ · rho vs log expert time $+0.379$ ($p = 0.056$)



Terminal-Bench's own long-horizon axis: expert estimate at or above four hours.